Simple TWAMP (STAMP) Extensions for Direct Measurement

draft-gandhi-ippm-stamp-direct-00

Rakesh Gandhi - Cisco Systems (rgandhi@cisco.com) - Presenter
Clarence Filsfils - Cisco Systems (cfilsfil@cisco.com)
Daniel Voyer - Bell Canada (daniel.voyer@bell.ca)
Mach(Guoyi) Chen - Huawei (mach.chen@huawei.com)
Bart Janssens - Colt (Bart.Janssens@colt.net)

Agenda

- Requirements and Scope
- Summary of Extensions
- Next Steps

Requirements and Scope

Requirements:

- Support stand-alone Direct Measurement for Data Packet Loss
- Very high scale for number of sessions and faster detection interval
  - Support hardware implementation

Goals:

- Avoid per session provisioning on Session-Reflector
- Avoid control-channel signaling for sessions

Scope:

- STAMP [RFC 8762]
- STAMP Extensions [RFC8972]

Stand-alone Direct Measurement Test Packet for Packet Loss

- Stand-alone Direct Measurement test packet defined
  - Hardware efficient counter-stamping
    - Well-known locations for transmit and receive traffic counters
  - Block number of the counters for alternate marking method [RFC 8321]
  - Traffic class of the counters for per class packet loss
- Direct Measurement test packet is also defined for authenticated mode
- User-configured destination UDP **Port2** is used for identifying direct measurement test packets
- Does not modify the existing STAMP procedure as different destination UDP port is used for direct measurement test packets
  - Other than Timestamp vs. Counter in the test packet, the protocol is same as STAMP
- Sequence Numbers allow to detect test packet loss, and connectivity loss
- Flags
  - X set to 1 for 64-Bit Counter, set to 0 for 32-Bit Counter
  - B set to 1 for Byte Counter, set to 0 for Packet Counter
  - T set to 1 for Sender-DSCP scoped Counter

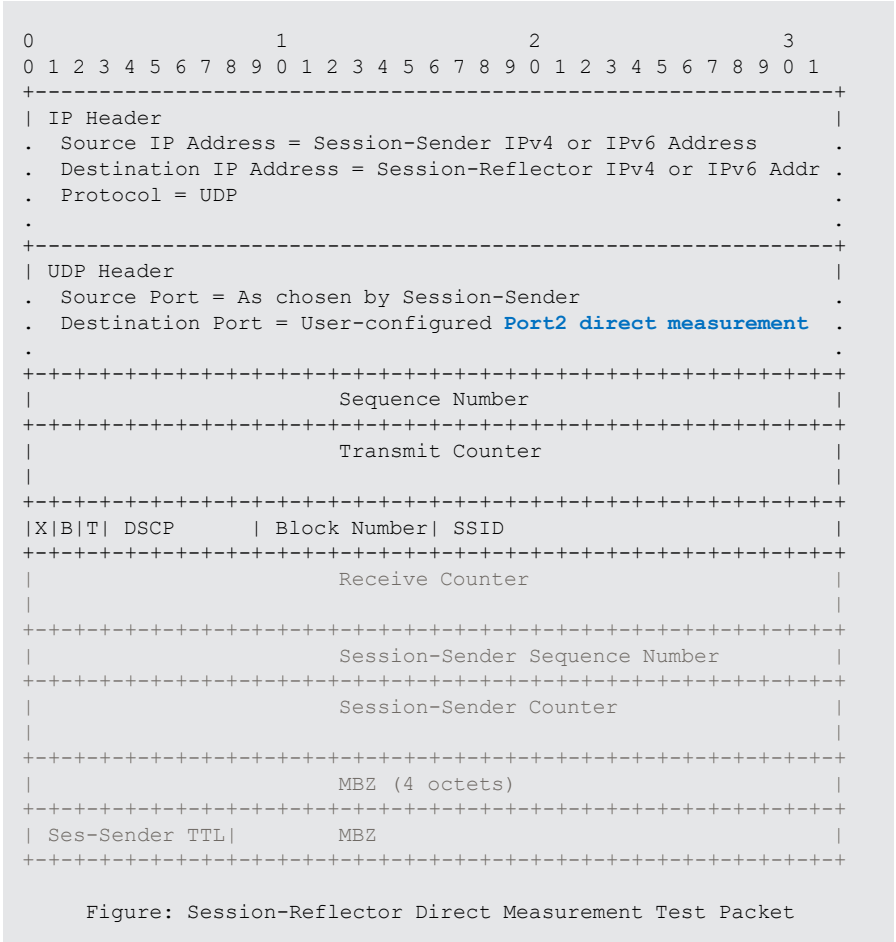

Figure: Session-Reflector Direct Measurement Test Packet

Direct Measurement TLV vs. Direct Measurement Test Packet

	Direct Measurement TLV	Direct Measurement Test Packet
TLV processing for each received packet on session-reflector in hardware	Yes	No
Need to write timestamp	Yes	No
Minimum bytes to load in write-able memory	114 (Eth 18, IPv6 40, UDP 8, STAMP 44, TLV Type 4, Total = 114 Byte)	70 (Eth 18, IPv6 40, UDP 8, Seq 4, Total = 70 Byte)
Counters at fixed location in the test packet for hardware counter-stamping	No (TLV-based)	Yes
Reply test packets with counters at the fixed location for hardware counter-stamping	No	Yes
Byte counters	No	Yes
64-bit packet and byte counters	No	Yes
Alternate-marking method packet loss - using block number for counters (out-of-order data packet support)	No	Yes
Per Traffic Class Counters	No	Yes

Link Loss Direct Measurement (P2P Circuits) - Inline Counter-stamping in Hardware

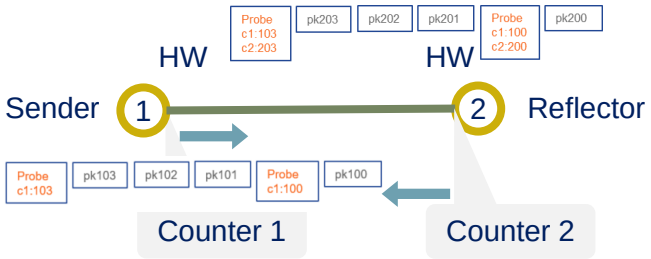

- Advertise extended TE metrics – link loss percentage
  - RFC 8570 (IS-IS)
  - RFC 7471 (OSPF)
  - RFC 8571 (BGP-LS)

- One Way Packet Loss %
  = 100* (((C1(t) – C1(t-1)) – (C2(t) – C2(t-1)) / (C1(t) – C1(t-1))
  = 100* (((103 – 100) – (203 – 200)) / (103 – 100))
  = 0
- Traffic Counters – counter-stamping in hardware

```
 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+
|   Type        |     Length    |
+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+
|A|  RESERVED   |                     Link Loss                   |
+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+
```

Alternate Marking Method for Packet Loss

- RFC 8321 - Alternate-Marking Method for Passive and Hybrid Performance Monitoring
- RFC 8957 - Synonymous Flow Label Framework
- Control plane-based packet loss with distributed forwarding LCs, using block number of the counters

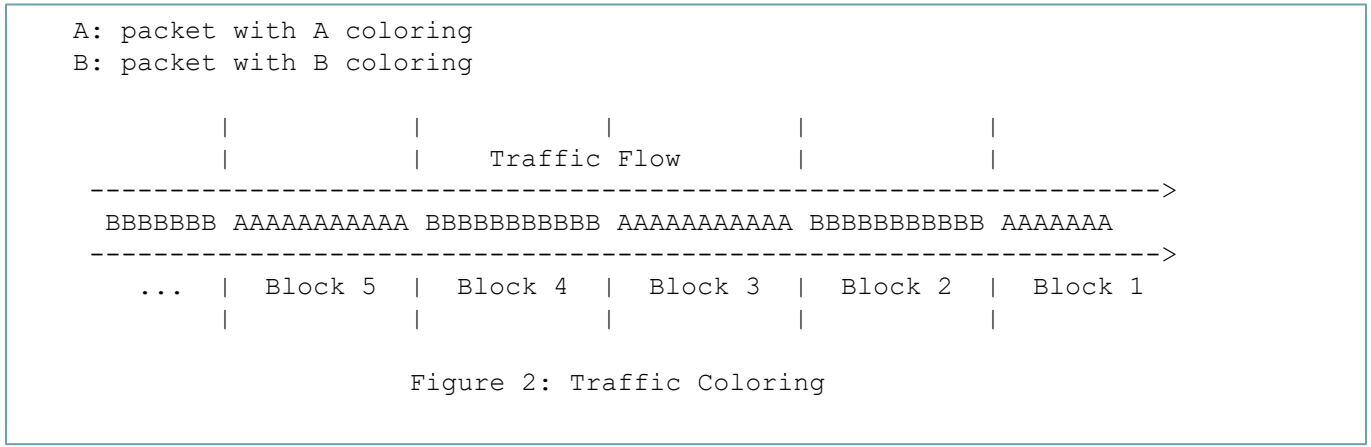

Figure 2: Traffic Coloring

Packet Loss Calculation

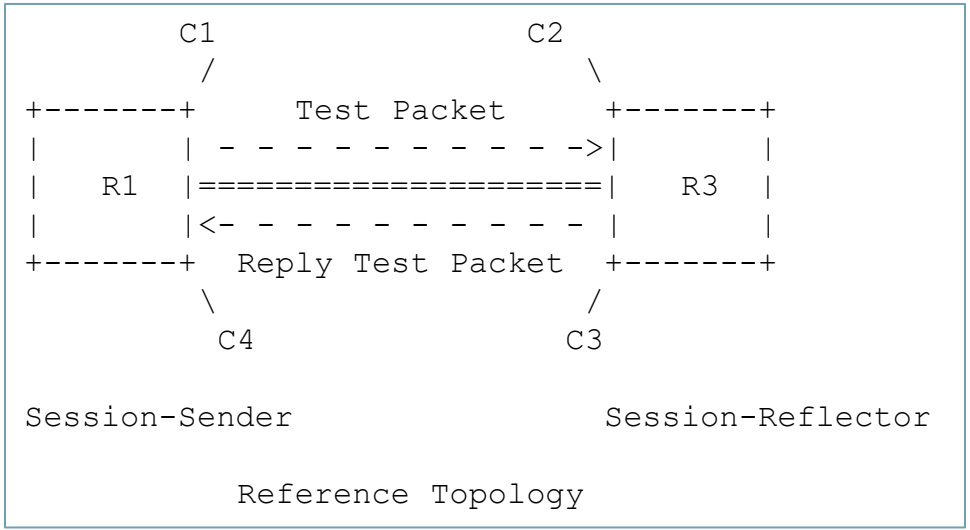

- Using the Counters C1, C2, C3 and C4 as per reference topology, from the n^{th} and $(n-1)^{th}$ direct measurement test packets.
- One-way receive packet loss[n-1, n] = (C2[n] - C2[n-1]) - (C1[n] - C1[n-1])
- Two-way receive packet loss[n-1, n] = (C4[n] - C4[n-1]) - (C3[n] - C3[n-1])
  + (C2[n] - C2[n-1]) - (C1[n] - C1[n-1])

Next Steps

- Welcome your comments and suggestions
- Requesting WG adoption

Thank you

STAMP Test Packet with Direct Measurement TLV

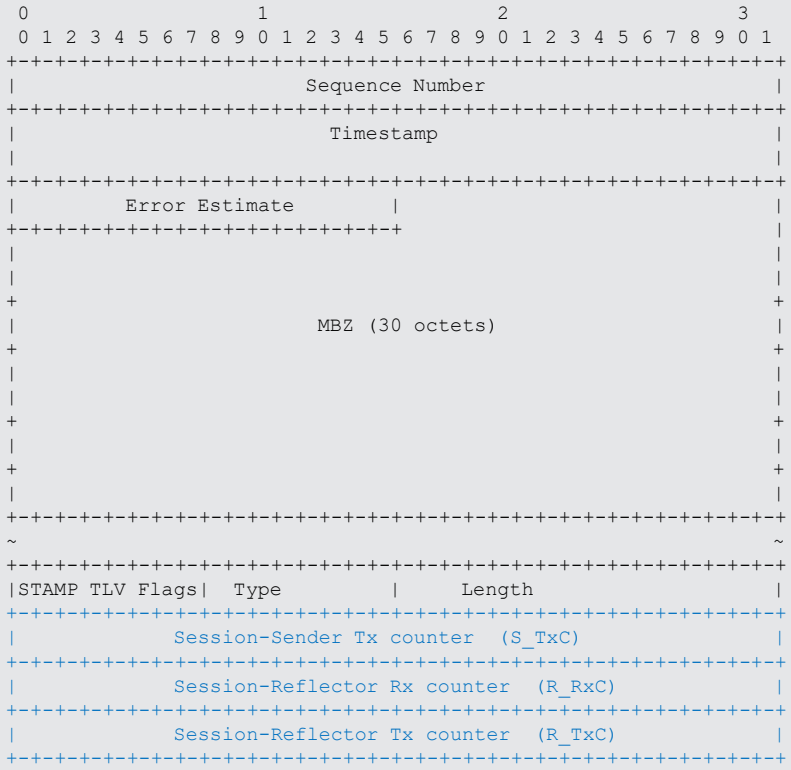

Figure: Session-Sender Test Packet Format

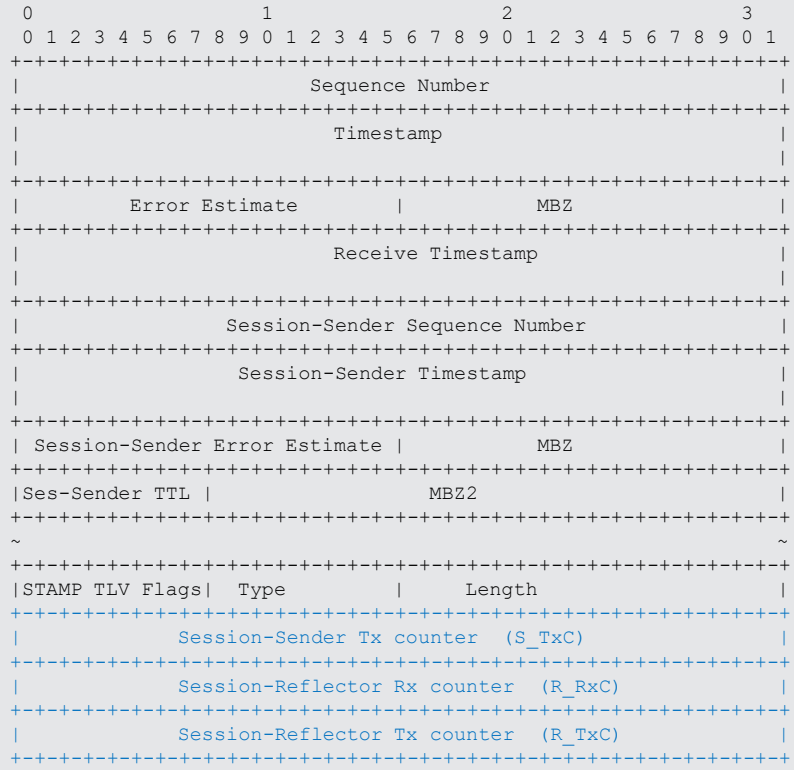

Figure: Session-Reflector Test Packet Format