

Performance Measurement Using TWAMP Light for Segment Routing Networks

draft-gandhi-spring-twamp-srpm-09

Rakesh Gandhi - Cisco Systems (rgandhi@cisco.com) - Presenter

Clarence Filsfils - Cisco Systems (cfilsfil@cisco.com)

Daniel Voyer - Bell Canada (daniel.voyer@bell.ca)

Mach(Guoyi) Chen - Huawei (mach.chen@huawei.com)

Bart Janssens - Colt (Bart.Janssens@colt.net)

Agenda

- Requirements and Scope
- History of the Draft
- Updates Since IETF-106
- Summary
- Next Steps

Requirements and Scope

Requirements:

- Delay and Loss Performance Measurement (PM)
 - ✓ Links and End-to-end P2P/P2MP SR Paths
 - ✓ Applicable to SR-MPLS/SRv6 data planes
- No need to negotiate UDP port to bootstrap PM session - spirit of SR
 - ✓ Stateless on egress node - spirit of SR
- Handle ECMP for SR Paths
- Support stand-alone direct-mode loss measurement

Scope:

- RFC 5357 (TWAMP Light) defined probe messages
- User-configured IP/UDP path for probe messages

History of the Draft

- Feb 2019
 - Draft was published - *draft-gandhi-spring-twamp-srpm-00*
- Mar 2019
 - Presented *draft-gandhi-spring-twamp-srpm-00* at IETF 104 Prague in SPRING WG
- May 2019
 - Added STAMP TLV for Return Path
- July 2019
 - Presented *draft-gandhi-spring-twamp-srpm-01* at IETF 105 Montreal in IPPM WG
 - Slide 9 Titled - Applicability of STAMP – STAMP is supported
- Aug 2019
 - *draft-gandhi-spring-twamp-srpm-02* included stand-alone LM messages
- Nov 2019
 - SPRING Chairs announced in the meeting the agreement with IPPM chairs to progress the draft in SPRING WG
 - Presented *draft-gandhi-spring-twamp-srpm-04* at IETF 106 Singapore in SPRING WG
- Mar 2020
 - Moved STAMP support to *draft-gandhi--spring-**stamp**-srpm-00*
 - Keep TWAMP Light support as informational in *draft-gandhi-spring-**twamp**-srpm-08*

Updates Since IETF-106 (Version-04)

Updates:

1. Defined Control Code for “In-band Response Requested” for TWAMP Light
 - ✓ Updated Two-way mode procedure using the Control Code
2. Moved STAMP support to a new draft - *draft-gandhi-spring-stamp-srpm-00*
3. Informational draft - as TWAMP Light is informational, see Appendix I in RFC 5357 and Appendix A RFC 8545
4. Moved Loopback mode to a new draft - *draft-gandhi-spring-sr-enhanced-plm-00*
5. Various editorial changes

Open Items:

- None

TWAMP Light Control Code Field

For a Query: Sender Control Code

0x0: Out-of-band Response Requested.

This is also the default behavior.

0x1: In-band Response Requested.

Indicates that this query has been sent over a bidirectional path and the probe response is required over the same path in the reverse direction.

The bidirectional path does not have to be an SR path.




Figure: Control Code in TWAMP Light Query Message

For a Response: Reflector Control Code

0x1: Error - Invalid Message.

Indicates that the operation failed because the received query message could not be processed.

0xN: Additional Error will be defined in future




Figure: Control Code in TWAMP Light Response Message

Performance Measurement Modes

- One-way Measurement Mode
 - Reply sent “out of band” on IP/UDP path - default
- Two-way Measurement Mode
 - Reply sent “in-band” on reverse SR path
 - Based on Control Code from the probe query message


Next Steps

- Welcome your comments and suggestions
- Implementation exists
- In WG adoption (SPRING WG) queue
- Keep IPPM WG in the loop about the milestones

Thank you

Backup

Example Provisioning Model



Probe Query for Links

- User-configured destination UDP **port1** is used for DM probe messages and **port2** is used for LM probe messages, both in unauthenticated mode.

```
+-----+
| IP Header                                     |
| . Source IP Address = Sender IPv4 or IPv6 Address . |
| . Destination IP Address = Reflector IPv4 or IPv6 Address . |
| . Protocol = UDP . |
| . . |
+-----+
| UDP Header                                   |
| . Source Port = As chosen by Sender . |
| . Destination Port = User-configured Port . |
| . . |
+-----+
| Payload = DM Message as specified in Section 4.2.1 of RFC 5357 |
| . Payload = DM Message as specified in Section 4.1.2 of RFC 5357. |
| . Payload = LM Message as specified in this document . |
| . . |
+-----+
```

Figure: Probe Query Message

Probe Query for SR-MPLS and SRv6 Policy

For **end-to-end** performance delay/loss measurement of SR Policy, the probe query message is sent on the SR Policy with:

1. MPLS label stack for SR-MPLS Policies
2. SRv6 SRH [RFC 8754] with SID list for SRv6 Policies

```

0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|                                     Segment(1)                        | TC  |S|          TTL          |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
.
.
.
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|                                     Segment(n)                        | TC  |S|          TTL          |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
| DM or LM Query Message including IP/UDP Header                    |
.
+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```

Figure: Example Probe Query Message for SR-MPLS Policy

```

0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| IP Header                                                                 |
. Source IP Address = Sender IPv6 Address                               .
. Destination IP Address = Next Hop IPv6 Address                       .
.                                                                       .
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| SRH as specified in RFC 8754                                           |
. <SID List>                                                             .
.                                                                       .
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| IP Header (Optional)                                                  |
. Source IP Address = Sender IPv6 Address                               .
. Destination IP Address = Reflector IPv6 Address                       .
.                                                                       .
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| UDP Header                                                             |
. Source Port = As chosen by Sender                                     .
. Destination Port = User-configured Port                               .
.                                                                       .
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
| Payload = DM or LM Query Message                                       |
.
+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```

Figure: Example Probe Query Message for SRv6 Policy

Probe Response Message

- The probe response message is sent using the IP/UDP information from the probe query message.

```
+-----+
| IP Header |
| . Source IP Address = Reflector IPv4 or IPv6 Address |
| . Destination IP Address = Source IP Address from Query |
| . Protocol = UDP |
| . |
+-----+
| UDP Header |
| . Source Port = As chosen by Reflector |
| . Destination Port = Source Port from Query |
| . |
+-----+
| Payload = DM Message specified in Section 4.2.1 of RFC 5357 | |
| . Payload = LM Message specified in this document |
| . |
+-----+
```

Figure: Probe Response Message

Stand-alone LM Message Format for TWAMP Light

- Loss Measurement (LM) message defined
 - Hardware efficient counter-stamping
 - Well-known locations for transmit and receive traffic counters
 - Stand-alone LM message, not tied to DM
- LM message format is also defined for authenticated mode
- User-configured destination UDP **port2** is used for identifying LM probe packets
- **Does not modify existing TWAMP Light (which is for DM) procedure as different UDP destination port2 is used for LM**

```
+-----+
| IP Header |
. Source IP Address = Sender IPv4 or IPv6 Address .
. Destination IP Address = Reflector IPv4 or IPv6 Address .
. Protocol = UDP .
. Router Alert Option Not Set .
.
+-----+
| UDP Header |
. Source Port = As chosen by Sender .
. Destination Port = User-configured Port2 for Loss Measurement. .
.
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|                               Sequence Number                       |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|                               Transmit Counter                       |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|X|B| Reserved | Block Number | MBZ |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|                               Receive Counter                       |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|                               Sender Sequence Number                 |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|                               Sender Counter                         |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|X|B| Reserved | Sender Block Nu| MBZ |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
| Sender TTL | Padding (3 Bytes) |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|                               Padding                                |
.
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
```

ECMP Support for SR Path

- SR Path can have ECMP between the ingress and transit nodes, between transit nodes and between transit and egress nodes.
- Sending PM probe queries that can take advantage of the hashing function in forwarding plane.
- Existing forwarding mechanisms are applicable to PM probe messages. Examples are:
 - For IPv4
 - Sweeping destination address in IPv4 header (e.g. 127/8)
 - For IPv6
 - Sweeping flow label in IPv6 header

Thank you